

EECS E6690 hw3

Chong Hu ch3467

Sep 22, 2019

P1

(a)

$$\begin{aligned}\mathbb{P}[Y = \text{true}] &= \frac{\exp(\hat{\beta}_0 + \hat{\beta}_1 X_1 + \hat{\beta}_2 X_2)}{1 + \exp(\hat{\beta}_0 + \hat{\beta}_1 X_1 + \hat{\beta}_2 X_2)} \\ &= \frac{\exp(-6 + 0.05 * 40 + 1 * 3.5)}{1 + \exp(-6 + 0.05 * 40 + 1 * 3.5)} \\ &= 0.3775407\end{aligned}$$

Hence, $\mathbb{P}[Y = \text{true}] = 0.3775407$.

(b)

$$\begin{aligned}\mathbb{P}[Y = \text{true}] &= \frac{\exp(\hat{\beta}_0 + \hat{\beta}_1 X_1 + \hat{\beta}_2 X_2)}{1 + \exp(\hat{\beta}_0 + \hat{\beta}_1 X_1 + \hat{\beta}_2 X_2)} = 0.5 \\ \exp(\hat{\beta}_0 + \hat{\beta}_1 X_1 + \hat{\beta}_2 X_2) &= 1 \\ -6 + 0.05 * x_1 + 1 * 3.5 &= 0 \\ x_1 &= 50\end{aligned}$$

The student in part (a) need to study 50 hours to have a 50% chance of getting an A in the class.

P2

$$\begin{aligned}f_{\text{Yes}}(x) &= \frac{1}{\sqrt{2\pi\hat{\sigma}^2}} \exp - \frac{(x - \mu_{\text{Yes}})^2}{2\hat{\sigma}^2} \\ f_{\text{No}}(x) &= \frac{1}{\sqrt{2\pi\hat{\sigma}^2}} \exp - \frac{(x - \mu_{\text{No}})^2}{2\hat{\sigma}^2} \\ \mathbb{P}[Y = \text{Yes} \mid X = 4] &= \frac{\pi_{\text{Yes}} f_{\text{Yes}}(x)}{\pi_{\text{Yes}} f_{\text{Yes}}(x) + \pi_{\text{No}} f_{\text{No}}(x)} \\ &= \frac{0.8 * 0.04032845}{0.8 * 0.04032845 + 0.2 * 0.05324133} \\ &= 0.7518524\end{aligned}$$

P3

$$\begin{aligned}\text{likelihood} &= \prod_{i=1}^n \left(\frac{\exp(\beta_0 + \beta_1 x_i)}{1 + \exp(\beta_0 + \beta_1 x_i)} \right)^{y_i} \left(\frac{1}{1 + \exp(\beta_0 + \beta_1 x_i)} \right)^{1-y_i} \\ l(\beta_0, \beta_1) &= \sum_{i=1}^n y_i(\beta_0 + \beta_1 x_i) - \ln(1 + \exp(\beta_0 + \beta_1 x_i))\end{aligned}$$

To maximize the likelihood, we need to set

$$\frac{\partial l(\beta_0, \beta_1)}{\partial \beta_0} = \sum_{i=1}^n \left(y_i - \frac{\exp(\hat{\beta}_0 + \hat{\beta}_1 x_i)}{1 + \exp(\hat{\beta}_0 + \hat{\beta}_1 x_i)} \right) = 0$$

$$\frac{\partial l(\beta_0, \beta_1)}{\partial \beta_1} = \sum_{i=1}^n x_i \left(y_i - \frac{\exp(\hat{\beta}_0 + \hat{\beta}_1 x_i)}{1 + \exp(\hat{\beta}_0 + \hat{\beta}_1 x_i)} \right) = 0$$

Using the Newton–Raphson algorithm, I write all variables in matrix form. \mathbf{W} is $n \times n$ diagonal matrix of weights with i th diagonal element $p(x; \beta^{\text{old}})(1 - p(x; \beta^{\text{old}}))$.

$$\mathbf{z} = \mathbf{X}\beta^{\text{old}} + \mathbf{W}^{-1}(\mathbf{y} - \mathbf{p})$$

$$\beta^{\text{new}} = (\mathbf{X}^T \mathbf{W} \mathbf{X})^{-1} \mathbf{X}^T \mathbf{W} \mathbf{z}$$

Then we start with $\beta_0 = 0$ and $\beta_1 = 0$ and perform 10 iterations.

```
library(matlib)

x = c(0.0, 0.2, 0.4, 0.6, 0.8, 1.0)
X = matrix(cbind(1, x), ncol = 2)
Y = matrix(c(0, 0, 0, 1, 0, 1), ncol = 1)
beta_old = matrix(c(0, 0), ncol = 1)
beta_new = beta_old
for(i in 1:10){
  p = exp(X %*% beta_old) / (1 + exp(X %*% beta_old))
  w = as.vector(p * (1 - p))
  W = diag(w)
  z = X %*% beta_old + inv(W) %*% (Y - p)
  beta_new = inv(t(X) %*% W %*% X) %*% t(X) %*% W %*% z
  beta_old = beta_new
}
beta_new

##           [,1]
## [1,] -4.097970
## [2,]  5.723309
```

After 10 iterations, we can get $\hat{\beta}_0 = -4.09797$ and $\hat{\beta}_1 = 5.7233085$.

P4

P5

P6

P7

```
library(ISLR)
library(tree)
attach(OJ)
set.seed(1000)
```

(a)

```
train = sample(dim(OJ)[1], 800)
OJ.train = OJ[train, ]
OJ.test = OJ[-train, ]
```

(b)

```
mod.tr <- tree(Purchase ~ ., data = OJ.train)
summary(mod.tr)
```

```
##
## Classification tree:
## tree(formula = Purchase ~ ., data = OJ.train)
## Variables actually used in tree construction:
## [1] "LoyalCH"      "PriceDiff"    "SalePriceMM"
## Number of terminal nodes: 8
## Residual mean deviance: 0.7486 = 592.9 / 792
## Misclassification error rate: 0.16 = 128 / 800
```

Training error rate for the tree is 0.16. The tree only use three variables LoyalCH, PriceDiff and SalePriceMM. There are 8 terminal nodes in the tree.

(c)

```
mod.tr

## node), split, n, deviance, yval, (yprob)
##      * denotes terminal node
##
## 1) root 800 1066.00 CH ( 0.61500 0.38500 )
##    2) LoyalCH < 0.5036 353 422.60 MM ( 0.28612 0.71388 )
##      4) LoyalCH < 0.276142 170 131.00 MM ( 0.12941 0.87059 )
##        8) LoyalCH < 0.035047 57 10.07 MM ( 0.01754 0.98246 ) *
##        9) LoyalCH > 0.035047 113 108.50 MM ( 0.18584 0.81416 ) *
##      5) LoyalCH > 0.276142 183 250.30 MM ( 0.43169 0.56831 )
##        10) PriceDiff < 0.05 78 79.16 MM ( 0.20513 0.79487 ) *
##        11) PriceDiff > 0.05 105 141.30 CH ( 0.60000 0.40000 ) *
##    3) LoyalCH > 0.5036 447 337.30 CH ( 0.87472 0.12528 )
##      6) LoyalCH < 0.764572 187 206.40 CH ( 0.75936 0.24064 )
##        12) SalePriceMM < 2.125 120 156.60 CH ( 0.64167 0.35833 )
##          24) PriceDiff < -0.35 16 17.99 MM ( 0.25000 0.75000 ) *
##          25) PriceDiff > -0.35 104 126.70 CH ( 0.70192 0.29808 ) *
##      13) SalePriceMM > 2.125 67 17.99 CH ( 0.97015 0.02985 ) *
##    7) LoyalCH > 0.764572 260 91.11 CH ( 0.95769 0.04231 ) *
```

Let's interpret the terminal node with "10)". The splitting variable of this node is PriceDiff. PriceDiff < 0.05 will be predicted as MM for the response Purchase. There are 78 points in total below this node. The deviation of those points is 79.16. 0.20513% points have CH as value for Purchase and 0.79487% points have MM as value for purchase. The * in the end means this is a terminal node.

(d)

```
plot(mod.tr)
text(mod.tr, pretty = 5)
```




Figure 1: Plot of the tree

Three split data points into two sets based on $\text{LoyalCH} < 0.5036$ at the root. Then, the subtrees split data points into different leaves based on different conditions. Among those conditions, we can clearly see that variable LoyalCH is used very often. In the end, we get eight terminal nodes to predict response **Purchase** as CH or MM.

(e)

```
mod.pred = predict(mod.tr, OJ.test, type = "class")
err_rate = sum(OJ.test$Purchase != mod.pred) / dim(OJ.test)[1]
print(sprintf("err rate: %.2f%% ", err_rate*100))
```

```
## [1] "err rate: 18.15% "
```

```
table(OJ.test$Purchase, mod.pred)
```

```
##      mod.pred
##      CH  MM
## CH 150  11
## MM  38  71
```

Error rate on test set is 18.15%.

(f)

```
tree.cv = cv.tree(mod.tr, FUN = prune.tree)
tree.cv
```

```
## $size
## [1] 8 7 6 5 4 3 2 1
##
## $dev
## [1] 663.1416 670.8785 669.0620 715.5108 737.4116 777.8179 770.9825
## [8] 1068.1253
##
## $k
## [1] -Inf 11.87503 12.41171 29.77434 31.80546 39.82936 41.38321
## [8] 306.37571
##
## $method
## [1] "deviance"
##
## attr("class")
## [1] "prune" "tree.sequence"
```

The optimal size is 8 based on deviation.

(g)

```
plot(tree.cv$size, tree.cv$dev, col = "blue", pch = 3,
     type = "b", xlab = "Tree Size", ylab = "cross-validated classification error")
```




Figure 2: cross-validated classification error vs. Tree size

(h)

When the tree size is 8, cross-validation error is the lowest.

(i)

```
tree.pruned = prune.tree(mod.tr, best = 8)
tree.pruned

## node), split, n, deviance, yval, (yprob)
##      * denotes terminal node
##
##  1) root 800 1066.00 CH ( 0.61500 0.38500 )
##    2) LoyalCH < 0.5036 353  422.60 MM ( 0.28612 0.71388 )
##      4) LoyalCH < 0.276142 170  131.00 MM ( 0.12941 0.87059 ) *
##        8) LoyalCH < 0.035047 57   10.07 MM ( 0.01754 0.98246 ) *
##        9) LoyalCH > 0.035047 113  108.50 MM ( 0.18584 0.81416 ) *
##      5) LoyalCH > 0.276142 183  250.30 MM ( 0.43169 0.56831 )
##        10) PriceDiff < 0.05 78   79.16 MM ( 0.20513 0.79487 ) *
##        11) PriceDiff > 0.05 105  141.30 CH ( 0.60000 0.40000 ) *
##    3) LoyalCH > 0.5036 447  337.30 CH ( 0.87472 0.12528 )
##      6) LoyalCH < 0.764572 187  206.40 CH ( 0.75936 0.24064 )
##        12) SalePriceMM < 2.125 120  156.60 CH ( 0.64167 0.35833 )
##          24) PriceDiff < -0.35 16   17.99 MM ( 0.25000 0.75000 ) *
##          25) PriceDiff > -0.35 104  126.70 CH ( 0.70192 0.29808 ) *
##        13) SalePriceMM > 2.125 67   17.99 CH ( 0.97015 0.02985 ) *
##      7) LoyalCH > 0.764572 260   91.11 CH ( 0.95769 0.04231 ) *

# create a pruned tree with five terminal nodes.
tree.pruned = prune.tree(mod.tr, best = 5)
tree.pruned

## node), split, n, deviance, yval, (yprob)
##      * denotes terminal node
##
##  1) root 800 1066.00 CH ( 0.61500 0.38500 )
##    2) LoyalCH < 0.5036 353  422.60 MM ( 0.28612 0.71388 )
##      4) LoyalCH < 0.276142 170  131.00 MM ( 0.12941 0.87059 ) *
##      5) LoyalCH > 0.276142 183  250.30 MM ( 0.43169 0.56831 ) *
##    3) LoyalCH > 0.5036 447  337.30 CH ( 0.87472 0.12528 )
##      6) LoyalCH < 0.764572 187  206.40 CH ( 0.75936 0.24064 )
##        12) SalePriceMM < 2.125 120  156.60 CH ( 0.64167 0.35833 ) *
##        13) SalePriceMM > 2.125 67   17.99 CH ( 0.97015 0.02985 ) *
##      7) LoyalCH > 0.764572 260   91.11 CH ( 0.95769 0.04231 ) *
```

(j)

```
summary(tree.pruned)

##
## Classification tree:
## snip.tree(tree = mod.tr, nodes = c(12L, 4L, 5L))
## Variables actually used in tree construction:
## [1] "LoyalCH"      "SalePriceMM"
## Number of terminal nodes:  5
## Residual mean deviance:  0.8138 = 646.9 / 795
## Misclassification error rate: 0.1962 = 157 / 800
```

The error rate of pruned tree on training set is 19.62% while the error rate of unpruned tree is 16%. After pruning, the error rate on training set is becoming higher. This is because we are removing some nodes in the tree and reducing the overfitting on training set.

(k)

```
pruned.pred = predict(tree.pruned, OJ.test, type = "class")
err_rate = sum(OJ.test$Purchase != pruned.pred) / dim(OJ.test)[1]
print(sprintf("err rate: %.2f%% ", err_rate*100))
```

```
## [1] "err rate: 21.11% "
```

The error rate of pruned tree on training set is 21.11% while the error rate of unpruned tree is 18.15%. In this scenario, the error rate of pruned tree is higher.